

✓ 1000-719bMSB Modeling of Complex Biological Systems

Deep Neural Network: Supervised Learning

✓ Basic python and pandas

<https://www.kaggle.com/lavanyashukla01/pandas-numpy-python-cheatsheet>

<https://www.utc.fr/~jlaforet/Suppl/python-cheatsheets.pdf>

List comprehensions are a concise way to create new lists from existing ones.

```
list1 = list(range(0,10))
print(list1)
```

```
→ [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

list1[0] # A vector in R starts with an index of 1. In Python, 0.

```
→ 0
```

```
list1[2:5]
```

```
→ [2, 3, 4]
```

```
list2 = []
for i in list1:
    list2.append(i+1)
```

```
print(list2)
```

```
→ [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
list3 = [i+1 for i in list1]
```

```
print(list3)
```

```
→ [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

✓ Classification of MNIST using densely connected layers

We are going to use the Keras library to implement a neural network that can classify handwritten digits - in just a few lines of code.

First we load and inspect the data. The dataset is split into training and test data.

```
import numpy as np
import tensorflow.keras as keras
import matplotlib.pyplot as plt
```

```
import tensorflow as tf
print(tf.__version__)
tf.compat.v1.disable_eager_execution()
```

```
→ 2.15.0
```

```
(train_images, train_labels), (test_images, test_labels) = keras.datasets.fashion_mnist.load_data()
```

```
→ Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 [=====] - 0s 0us/step
```

```
train_images.shape
```

```
(60000, 28, 28)
```

```
train_labels.shape
```

```
(60000,)
```

```
test_images.shape
```

```
(10000, 28, 28)
```


```
test_labels.shape
```

```
(10000,)
```

Let's plot one of the digits and the corresponding label.

```
print('Label of element 0:', train_labels[1])
plt.matshow(train_images[1], cmap='gray')
plt.show()
```

```
Label of element 0: 0
```



In this step we define the neural network. ReLu is an activation function defined as $f(x) = \max(0, x)$.

Softmax activation function is normalized such that the sum of all outputs is equal 1.

```
from tensorflow.keras import layers
from tensorflow.keras import models
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
model.add(layers.Dense(10, activation='softmax'))
```

With compile we tell the network which optimizer and loss function to use. Optimizer specifies the particular implementation of the gradient-descent, e.g. how it adapts the learning rate. 'Metrics' specifies the output during the training.

```
model.compile(optimizer='rmsprop',
              loss='mean_squared_error',
              metrics=['accuracy'])
```

```
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dense_5 (Dense)	(None, 10)	5130
Total params: 407050 (1.55 MB)		
Trainable params: 407050 (1.55 MB)		
Non-trainable params: 0 (0.00 Byte)		

We are using a densely connected network, so we have to flatten the images.

Input values should be in the range (0,1) for fast convergence.

```
train_images_flat = train_images.reshape((60000, 28 * 28))
train_images_flat = train_images_flat.astype('float32') / 255
test_images_flat = test_images.reshape((10000, 28 * 28))
test_images_flat = test_images_flat.astype('float32') / 255
```

Convert the labels to a 'one-hot' coding.

```
from tensorflow.keras.utils import to_categorical
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```


```
train_images.reshape((60000, 28*28)).shape
```

```
(60000, 784)
```

```
train_labels[0]
```

```
array([0., 0., 0., 0., 0., 0., 0., 0., 0., 1.], dtype=float32)
```

```
plt.matshow(train_images[0], cmap='gray')
plt.show()
```



```
model.fit(train_images_flat, train_labels, epochs=5, batch_size=128)
```

```
Train on 60000 samples
Epoch 1/5
60000/60000 [=====] - 4s 64us/sample - loss: 0.0311 - accuracy: 0.7848
Epoch 2/5
60000/60000 [=====] - 4s 61us/sample - loss: 0.0208 - accuracy: 0.8569
Epoch 3/5
60000/60000 [=====] - 5s 78us/sample - loss: 0.0184 - accuracy: 0.8721
Epoch 4/5
60000/60000 [=====] - 4s 68us/sample - loss: 0.0172 - accuracy: 0.8828
Epoch 5/5
60000/60000 [=====] - 4s 62us/sample - loss: 0.0163 - accuracy: 0.8888
<keras.src.callbacks.History at 0x7b50944ebbe0>
```

Let's check the performance on the test set. If the accuracy is less than the training accuracy, then we might be overfitting!

```
test_loss, test_acc = model.evaluate(test_images_flat, test_labels)
print('test_acc:', test_acc)
```


```
test_acc: 0.8698
```

We can also find the predictions for a selection of input images.

```
predictions = model.predict(train_images_flat[:10])
```

```
img_num = 3
print(predictions[img_num])
print(train_labels[img_num])
plt.matshow(train_images[img_num], cmap='gray')
plt.show()
```

```
[6.0702302e-02 2.2748517e-04 5.0025887e-05 8.9770937e-01 3.4096718e-06
 7.5253497e-06 4.1277580e-02 4.1745642e-07 2.0891413e-05 1.0804190e-06]
[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
```



```
train_labels
```

```
array([[0., 0., 0., ..., 0., 0., 1.],
       [1., 0., 0., ..., 0., 0., 0.],
       [1., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [1., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

✓ Classification of MNIST using convolutional layers

We have build a classifier for handwritten images only using densely connected layers. Let's see if we can do better using convolutional layers!

First define the convolutional layers.

```
model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model2.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 64)	36928
Total params: 55744 (217.75 KB)		
Trainable params: 55744 (217.75 KB)		
Non-trainable params: 0 (0.00 Byte)		

Now add a classifier on top of the convnet.

```
model2.add(layers.Flatten())
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dense(10, activation='softmax'))
```

```
model2.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_6 (Dense)	(None, 64)	36928
dense_7 (Dense)	(None, 10)	650

=====
Total params: 93322 (364.54 KB)
Trainable params: 93322 (364.54 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```
train_images_conv = train_images.reshape((60000, 28, 28, 1))
train_images_conv = train_images_conv.astype('float32') / 255
test_images_conv = test_images.reshape((10000, 28, 28, 1))
test_images_conv = test_images_conv.astype('float32') / 255
```

```
model2.compile(optimizer='rmsprop',
loss='categorical_crossentropy',
metrics=['accuracy'])
```

```
model2.fit(train_images_conv, train_labels, epochs=15, batch_size=64)
```

Train on 60000 samples

Epoch 1/15
60000/60000 [=====] - 45s 743us/sample - loss: 0.5382 - accuracy: 0.8007
Epoch 2/15
60000/60000 [=====] - 45s 758us/sample - loss: 0.3292 - accuracy: 0.8791
Epoch 3/15
60000/60000 [=====] - 44s 739us/sample - loss: 0.2777 - accuracy: 0.8975
Epoch 4/15
60000/60000 [=====] - 44s 729us/sample - loss: 0.2470 - accuracy: 0.9096
Epoch 5/15
60000/60000 [=====] - 44s 729us/sample - loss: 0.2240 - accuracy: 0.9171
Epoch 6/15
60000/60000 [=====] - 45s 755us/sample - loss: 0.2068 - accuracy: 0.9244
Epoch 7/15
60000/60000 [=====] - 44s 732us/sample - loss: 0.1893 - accuracy: 0.9309
Epoch 8/15
60000/60000 [=====] - 44s 731us/sample - loss: 0.1752 - accuracy: 0.9356
Epoch 9/15
60000/60000 [=====] - 44s 735us/sample - loss: 0.1636 - accuracy: 0.9391
Epoch 10/15
60000/60000 [=====] - 45s 749us/sample - loss: 0.1494 - accuracy: 0.9455
Epoch 11/15
60000/60000 [=====] - 44s 731us/sample - loss: 0.1394 - accuracy: 0.9482
Epoch 12/15
60000/60000 [=====] - 43s 724us/sample - loss: 0.1308 - accuracy: 0.9509
Epoch 13/15
60000/60000 [=====] - 45s 747us/sample - loss: 0.1203 - accuracy: 0.9551
Epoch 14/15
60000/60000 [=====] - 44s 737us/sample - loss: 0.1153 - accuracy: 0.9571
Epoch 15/15
60000/60000 [=====] - 44s 727us/sample - loss: 0.1075 - accuracy: 0.9599
<keras.src.callbacks.History at 0x7b50a9797df0>

```
test_loss, test_acc = model2.evaluate(test_images_conv, test_labels)
print(test_acc)
```

0.9088

✓ Introducing Fashion MNIST (Homework dataset)

The MNIST dataset is not too demanding, let's try something a little more difficult - Fashion MNIST.

[LINK TO IMAGE](#)

Check out labels on [GitHub](#)

```
(train_imgs_fash, train_labels_fash), (test_imgs_fash, test_labels_fash) = keras.datasets.fashion_mnist.load_data()
```


```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
32768/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26427392/26421880 [=====] - 1s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
8192/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4423680/4422102 [=====] - 0s 0us/step
```

```
train_imgs_fash.shape
```

```
(60000, 28, 28)
```

```
plt.imshow(train_imgs_fash[12], cmap=plt.get_cmap('gray'))
```

```
<matplotlib.image.AxesImage at 0x7fd36038a850>
```



HOMEWORK 1

Build a classifier for fashion MNIST.

- 1. Use exactly the same architectures (both densely connected layers and from convolutional layers) as the above MNIST** e.g., replace the dataset. Save the Jupyter Notebook in its original format and output a PDF file after training, testing, and validation. Make sure to write down how do they perform (training accuracy, testing accuracy).
- 2. Improve the architecture.** Experiment with different numbers of layers, size of layers, number of filters, size of filters. You are required to make those adjustment to get the highest accuracy. Watch out for overfitting – we want the highest testing accuracy! Please provide a PDF file of the result, the best test accuracy and the architecture (different numbers of layers, size of layers, number of filters, size of filters)

✓ Visualizing Filter Response

We use gradient descent in input space to display the visual pattern each filter is maximally responsive to. To this end we take a VGG19 convnet pretrained on the ImageNet dataset.

[Very Deep Convolutional Networks for Large-Scale Image Recognition](#) Karen Simonyan, Andrew Zisserman

[DL Architecture](#)

```
from tensorflow.keras.applications import VGG19
from tensorflow.keras import backend as K
import numpy as np

import matplotlib.pyplot as plt

#Load pretrained model
#we omit the densely connected layers of the network
model = VGG19(weights='imagenet', include_top=False)

↗ Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_n
80134624/80134624 [=====] - 0s 0us/step
```

```
model.summary()
```

↗ Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_conv4 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_conv4 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_conv4 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0

=====

Total params: 20024384 (76.39 MB)
Trainable params: 20024384 (76.39 MB)
Non-trainable params: 0 (0.00 Byte)

```
#Specify filter you want to visualize and get its output
layer_name = 'block5_conv3'
filter_index = 3
layer_output = model.get_layer(layer_name).output

#Loss is the averaged activation of the chosen filter
loss = K.mean(layer_output[:, :, :, filter_index])
```

```
#Gradients of loss with respect to the input
#upgrading to 2.x: tf.gradients is no longer supported
#requiring tf.compat.v1.disable_eager_execution()
grads = K.gradients(loss, model.input)[0]

#A trick is to normalize the gradients by their L2 norm
#This ensures that the magnitude of the gradients is always in the same range
#and leads to a smooth descent process
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)

#The tensors defined so far (loss, grads) were symbolic
#To obtain values we need to feed an input via K.function

iterate = K.function([model.input], [loss, grads])
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])

print(grads)
print(grads_value)
```

```
↳ Tensor("truediv:0", shape=(None, None, None, 3), dtype=float32)
[[[[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  ...

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
  ...
  [0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]]]]
```

```
#Implement the actual gradient descent
#Initial input is a grey image with some noise

input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step

print(grads_value)
```

```
↳ [[[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]
```



```

...
[0. 0. 0.]
[0. 0. 0.]
[0. 0. 0.]]

[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]

...
[0. 0. 0.]
[0. 0. 0.]
[0. 0. 0.]]

[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]

...
[0. 0. 0.]
[0. 0. 0.]
[0. 0. 0.]]

...

[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]

...
[0. 0. 0.]
[0. 0. 0.]
[0. 0. 0.]]

[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]

...
[0. 0. 0.]
[0. 0. 0.]
[0. 0. 0.]]]]

```

#Postprocess to turn into displayable image

```

def deprocess_image(x):
    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1

    x += 0.5
    x = np.clip(x, 0, 1)

    x *= 255
    x = np.clip(x, 0, 255).astype('uint8')
    return x

plt.imshow(deprocess_image(input_img_data[0]))

```



<matplotlib.image.AxesImage at 0x7b509443f4c0>

